



Ordinary memory controls reproduce the recovery deficit


12,000 repair updates · all three matched schedules shown · black marks = means



One damaged parent; schedules are not independent pretrained models. Each point: 768 requests.
Recovery gates include per-cell and late-half criteria. No catastrophic cognition failure is established.

The failures are concentrated in particular tasks

Original validation panels · 256 requests per task and schedule · LN-113




Circle / square / triangle = schedules 1 / 2 / 3; black marks = means.

Task-specific failures do not imply general cognitive collapse; all replications are shown.

Lookup accuracy across memory loads



Diagnostic panels · thin lines = individual schedules · thick lines = means · LN-113



Length 2: only 2 unique cores and a 100% constant baseline. Length 4: only 29 unique cores.
256 requests per schedule and length; repeated/overlapping panels, not a new generalization test.

Useful bypasses fit inside the tiny machine budget


Exhaustive 4,096-transition panel · synthetic authorization · no learning · LN-136



Disclosure denominator: all 2,048 denied requests. Selective patch: 1,024/1,024 targeted successes.
Both edits fit one word; caller substitution adds 1 instruction on denied requests but remains within 59.

Changing the identity input can preserve useful execution

Explanatory diagram of one bypass family · actual external identities stay fixed



For the fixed circuit grammar: $\text{Step}[T(C)](q,s,c,o) = \text{Step}[C](q,s,o,o)$.

This is equivalence to transformed authorized histories under explicit premises, not a universal SCC no-go.